



Intro to Computer Vision



What is Computer Vision?

- Broadly, a field of computer science focused on understanding images and videos
- Subfields include:
 - Scene reconstruction
 - Object detection
 - Object recognition
 - Activity recognition
 - Pose estimation
 - Motion estimation
 - ...and much more!
- Where is it used?
 - Almost everywhere!
 - Healthcare: Medical imaging analysis, disease detection
 - Transportation: Self-driving vehicles, license plate recognition
 - Entertainment: Face filters, AR/VR



Imaging Basics

- A pixel: Smallest unit of a digital image
- Resolution: Number of pixels in an image (width x height)
- Color depth: Bits needed to represent each pixel(bpp)
 - 1 bit: (0-1)
 - 8-bit (0-255)
 - 24-bit (RGB)
 - 32-bit (RGBA, supports transparency)
- Color spaces:
 - RGB
 - BGR
 - Grayscale
- HSV(Bettering Human Perception)
 - Color type, Saturation, Value (Brightness)
- Most computers store images as **matrices**
 - 2D array for grayscale images
 - 3D array for color images
 - Each element represents one pixel



[107	102	98	255]	[89	88	81	255]	[84	83	76	255]
[171	161	151	255]	[171	161	155	255]	[192	184	179	255]
[111	109	100	255]	[104	105	96	255]	[102	103	94	255]
[115	114	101	255]	[115	116	105	255]	[115	117	107	255]
[156	150	148	255]	[156	150	148	255]	[165	160	157	255]
...											
[145	139	123	255]	[122	116	105	255]	[115	109	99	255]
[90	89	74	255]	[90	89	74	255]	[92	90	77	255]
[147	141	125	255]	[126	120	108	255]	[119	113	103	255]

Basic Image Processing

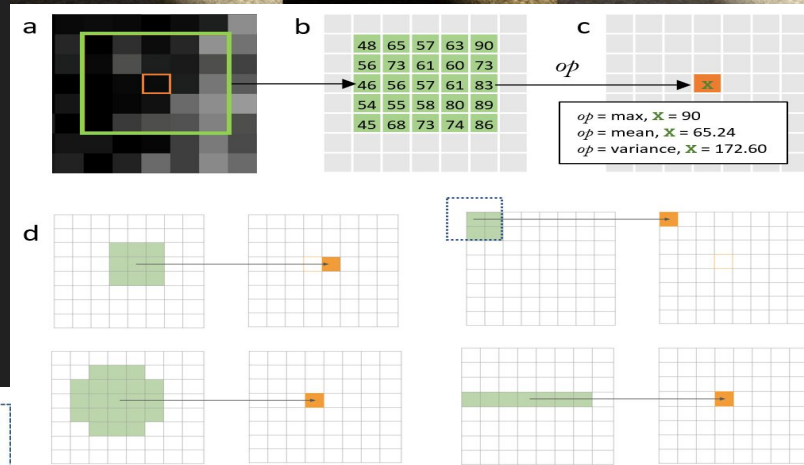
1. Point Operations

- Brightness adjustment (+ or - value)
- Contrast enhancement (multiplication)
- Thresholding (binary conversion)
- Gamma correction (power law)




2. Geometric Transformations

- Scaling (resize)
- Rotation (angle)
- Translation (move)
- Flipping (mirror)
- Affine transformations



3. Neighborhood Operations

- Smoothing (blur)
- Sharpening
- Edge detection
- Noise reduction
- Median filtering



What we will be using today

- **Google Colab:** For sharing code
- **Python:** Programming language commonly used for machine learning
- **OpenCV:** Library that provides functions for real-time and offline computer vision
- **PyTorch:** Deep learning library used for constructing neural networks
- **Torchvision:** PyTorch extension containing functions for image transformations and augmentations
- **Matplotlib:** For visualizing our data



Notebook With Exercises



<https://tinyurl.com/uclacvdatateach>

Exercise 1: Loading and basic image processing

- Since most image operations happen on a matrix representation of the image, our first step is to “read” the image as matrix
- We can do this with `cv2.imread(“/path/to/image”)`
- For efficiency, OpenCV converts an image to a **Numpy array**, which is more efficient than a Python list
 - You can get the data type of a Numpy array with `.dtype`
 - You can get the max and min with `.max()` and `.min()` respectively
- Since color images are represented by their red, green, and blue components, we can break apart images and “stitch” them back together




Exercise 2: Basic Image Operations


- **Resizing**
 - Involves scaling the dimensions of an image
 - It preserves the original proportions
- **Cropping**
 - Extracts a specific portion of the image
 - Removes unwanted outer areas
 - Helps focus on the main subject
- **Rotation**
 - Turns the image around its center point
 - Can be done in degrees (90° , 180° , 270°)
 - May require padding or cropping to handle corners
- **Flipping**
 - Mirrors the image horizontally or vertically
 - Useful for data augmentation
 - Helps models learn orientation invariance

- `cv2.resize(image, (new_width, new_height), interpolation)`
 - Takes in three parameters: image to resize, (width, height), interpolation

Counterclockwise rotation




Clockwise rotation



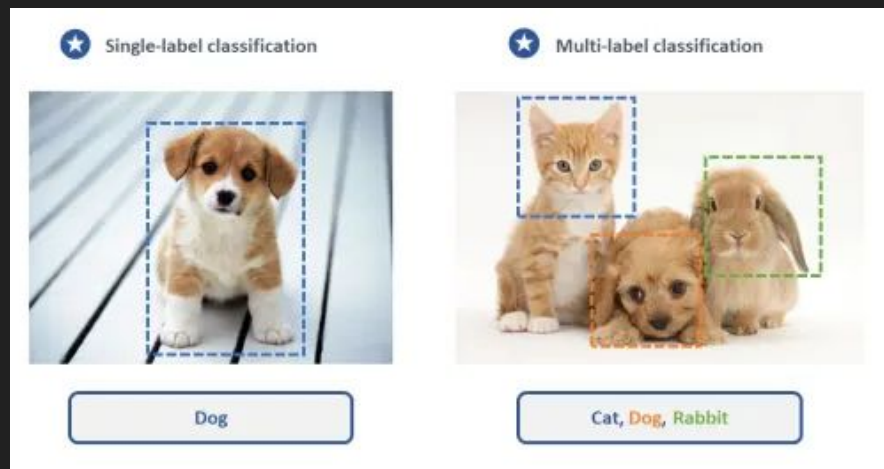
Exercise 3: Landmark Detection

- Gaussian Blur
 - Reduces image noise and detail
 - Uses a Gaussian kernel (bell-shaped) to “smooth” out an image
 - Parameters:
 - Size of the kernel
 - Kernel intensity
 - Essentially preprocessing
- Contour Drawing
 - Finds areas of intensity in the image
 - Find contours (cv2.findContours)
 - Draw contours (cv2.drawContours)
 - Applications
 - Object counting
 - Shape detection
 - Feature extraction
- Background Subtraction
 - Thresholding (cv2.threshold)
 - Creates a binary mask and applies it to the image
 - Used for: isolating objects, removing unwanted objects, preparing images for contour drawing




Machine Learning in Computer Vision

- Machine Learning: Process of teaching computers to learn from data without explicit programming
 - System learns patterns and features from examples instead of following rules
- Very useful for images, where data is often unstructured and is difficult to consider all scenarios
- ML can
 - Handle variations in lighting, pose, scale
 - Learn relevant features automatically
 - Generalize to new, unseen examples
 - Adapt to different conditions
- Types of Training data
 - Labeled images for supervised learning
 - Unlabeled images for unsupervised learning
- Use cases
 - Classification tasks (object recognition)
 - Detection tasks (bounding boxes)
 - Segmentation tasks (pixel-level labels)
 - Can require significant human annotation effort




Deep Learning in Computer Vision

- Deep Learning: subset of machine learning focused on **neural nets**
 - Key components:
 - Neurons (nodes)
 - Layers
 - Weights and biases
 - Activation functions (ReLU, sigmoid, tanh)
 - Loss functions
 - Optimizers (SGD, Adam, RMSprop)
- Convolutional Neural Network (CNN)
 - A convolutional neural network is specifically designed to process grid-like data (i.e. images)
 - Layers:
 - Convolutional Layers: extract features from images
 - Activation Function: adds non-linearity to our network
 - Pooling Layers: reduces spatial dimensions
 - Fully Connected Layers: final classification/regression



Deep Learning in Computer Vision

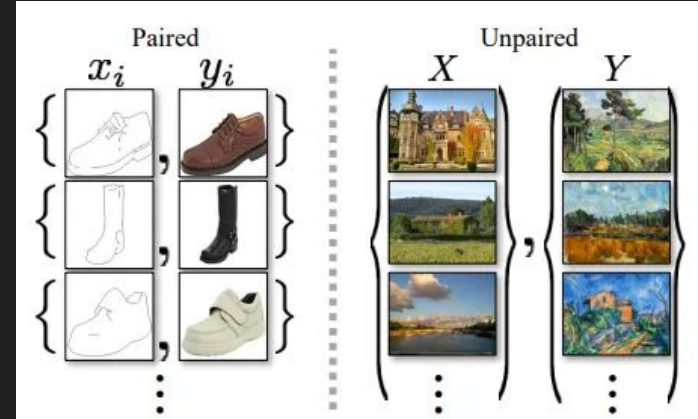
- Deep Learning can have some very cool applications in computer vision
- Image Classification
 - Identifying objects, photos, or categories in images
- Object Detection
 - Locating and identifying multiple objects in a single image
- Face Recognition
 - Identifying and verifying individuals from facial features
- Image Generation
 - Creating new images from descriptions or other images
- Pose Estimation
 - Detecting human body positions and movements
- Style Transfer
 - Applying artistic styles to regular photographs



Today, we will be creating a
style transfer model that
applies a cartoon artistic style!

We will be creating a model for **style transfer**!

- **Generative Adversarial Network:** A deep learning model used often for image generation
 - A GAN has two components: a Generator and a Discriminator
 - The Generator learns to generate images (in this case, transform images to a cartoon style) and the Discriminator tries to “guess” whether the generated images are real or fake
- We will be implementing **CartoonGAN** – a GAN specifically optimized for cartoons!
 - Instead of just predicting between fake and real images, CartoonGAN tries to predict against a smoothed version of the cartoon images, which makes it generate better images
 - CartoonGAN works on **unpaired** images - images that do not correspond to one another



Model Training

- **Step 1: Train Discriminator**

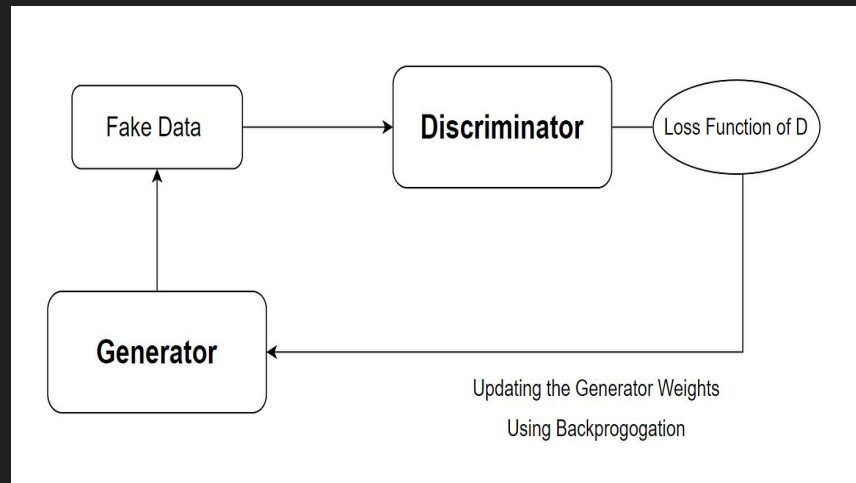
- Shows D three types of images:
 - Real cartoons (should say "real")
 - Smoothed cartoons (should say "fake")
 - Generated images (should say "fake")
- Like training an art critic to spot different styles

- **Step 2: Train Generator**

- G tries to create images that:
 - Look like cartoons (fool D)
 - Preserve original photo content
- Like an artist learning to balance style and content

- **Validation & Checkpointing**

- Regularly tests G's performance on new photos
- Saves the best performing model
- Like keeping your best artwork for reference



Model Inference

- At the end, our trained Generator can be used for generating realistic cartoon images!
- We will use our Generator specifically to convert images to cartoons